

The Hitchhiker's Guide to Fuzzer Coverage Metrics

The Hitchhiker's Guide to Fuzzer Coverage Metrics



Me as a grad student
(sometimes)

```
$ whoami
```

- > 10 years as a security researcher @ DSTG
- Principal vulnerability researcher @ Interrupt Labs
- Just submitted my PhD @ ANU 🎉

```
$ whoami
```

- > 10 years as a security researcher @ DSTG
- Principal vulnerability researcher @ Interrupt Labs
- Just submitted my PhD @ ANU 🎉



Me as a grad student
(sometimes)

An Aside, DSTG

- Government R&D
- Worked on all things “cyber”
 - Computer network defence
 - Malware analysis
 - Vulnerability research

Another Aside, Interrupt Labs

Vulnerability Research

Analyze a product to find, understand, or exploit one or more vulnerabilities

Lots of crossover with software testing

Last Aside, PhD

- Supervised by **Tony Hosking** (ANU) + **Mathias Payer** (EPFL)
- Quickly realised academia wasn't for me, but I had lots of transferable skills
 - Subject matter expertise
 - How to complete a nebulously-defined project
 - How to write
 - ...

Let's talk (more)
about fuzzing

What is Fuzzing?

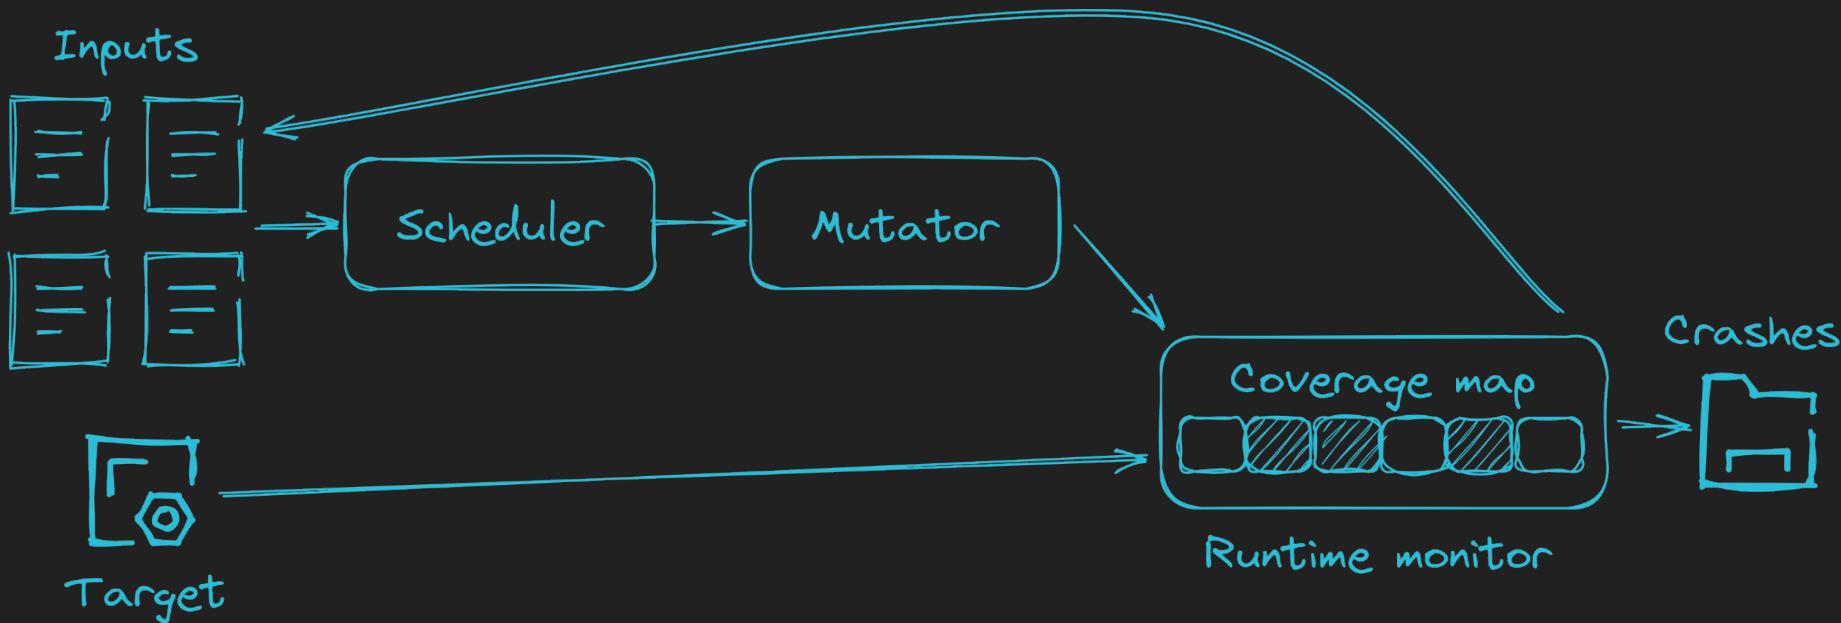
1. Generate random/invalid data
2. Execute target with said data
3. See if target breaks
4. Return to 1.

What is Fuzzing?

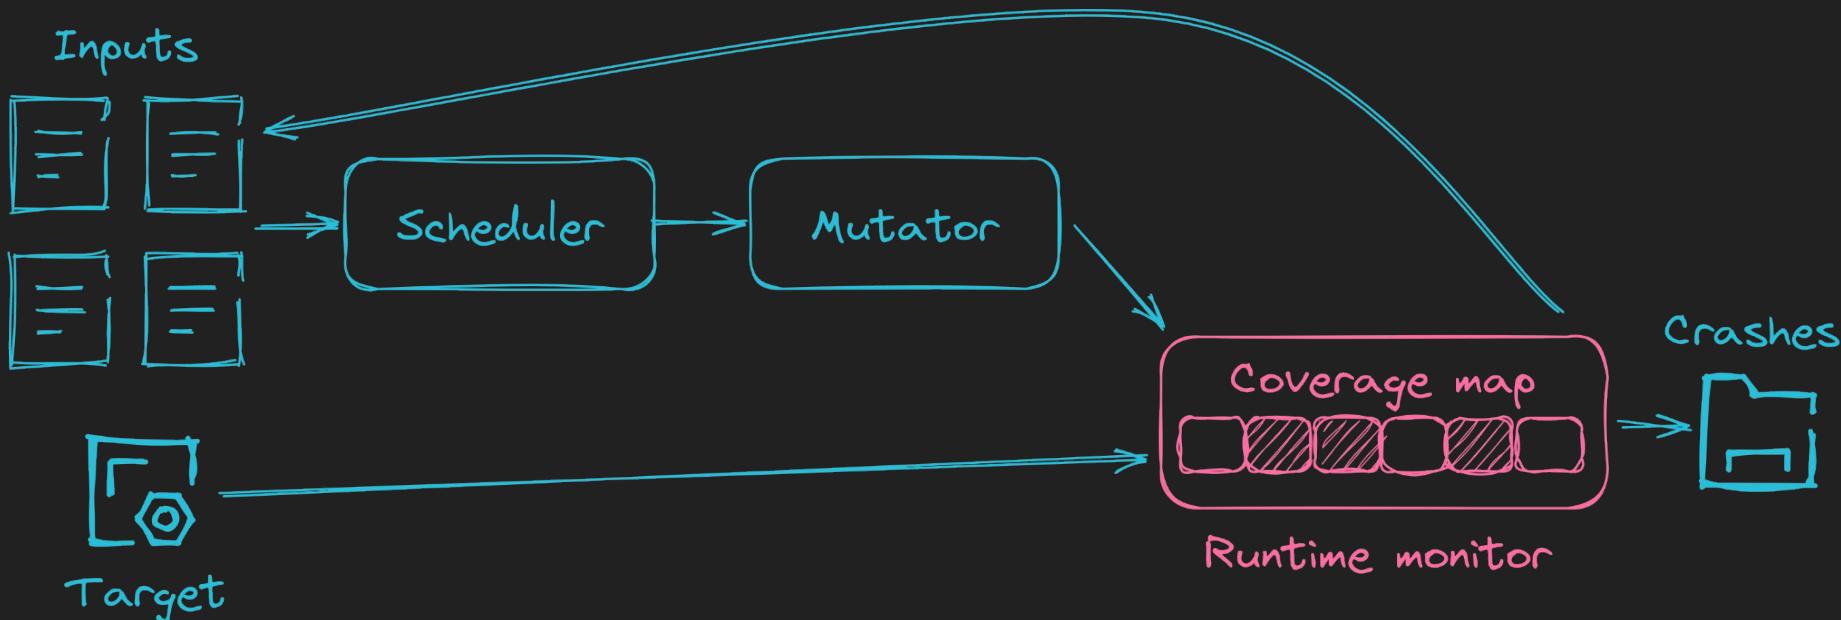
1. Generate random/invalid data
2. Execute target with said data
3. See if target breaks
4. Return to 1.

Maybe a bit more to it 🤔

Fuzzing 101



Fuzzing 101



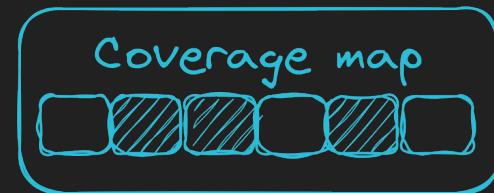
Fuzzing 101

Fuzzers find bugs by exploring a target's
state space

Approximate the target's state space and
track at runtime

- Must be lightweight!

Retain inputs uncovering new states



Runtime monitor

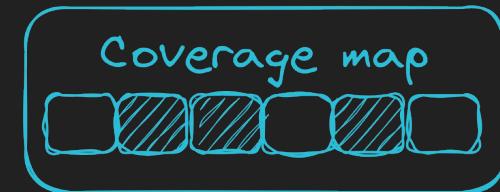
Fuzzing 101

Fuzzers find bugs by exploring a target's *state space*

Approximate the target's state space and track at runtime

- Must be lightweight!

Retain inputs uncovering new states



Runtime monitor

You can't find bugs in states never covered

How do we measure coverage?



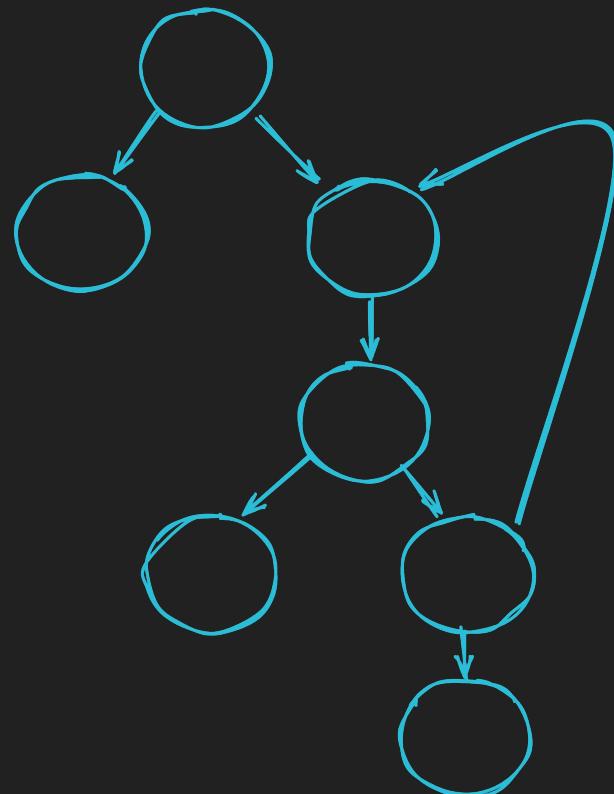
Abstraction!

Approximate program states

- Control flow
- Data flow

Control Flow

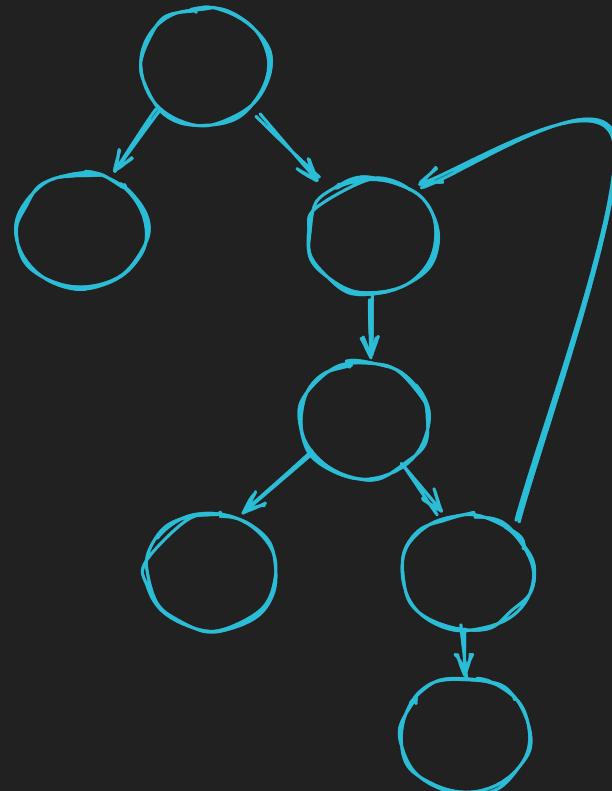
Decompose a function into a control-flow graph



Control Flow: Basic Blocks

Decompose a function into a
control-flow graph

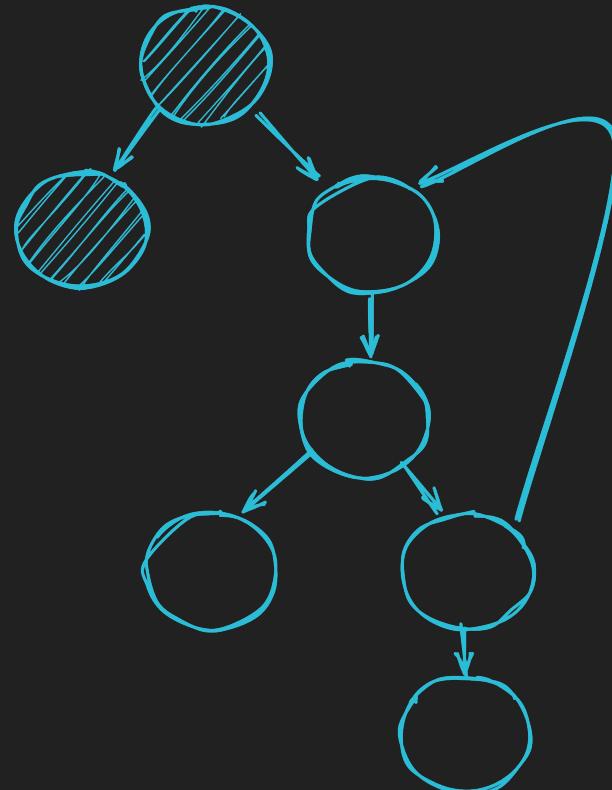
Record when nodes are covered



Control Flow: Basic Blocks

Decompose a function into a
control-flow graph

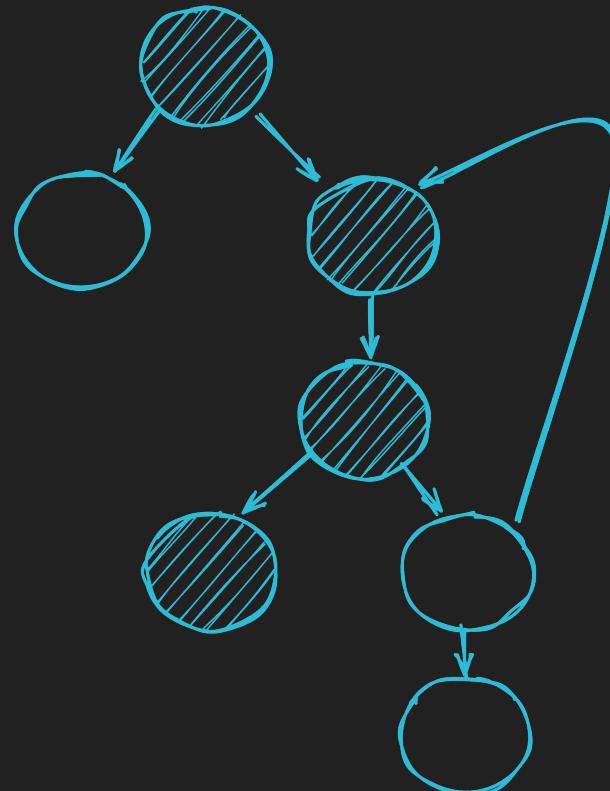
Record when nodes are covered



Control Flow: Basic Blocks

Decompose a function into a
control-flow graph

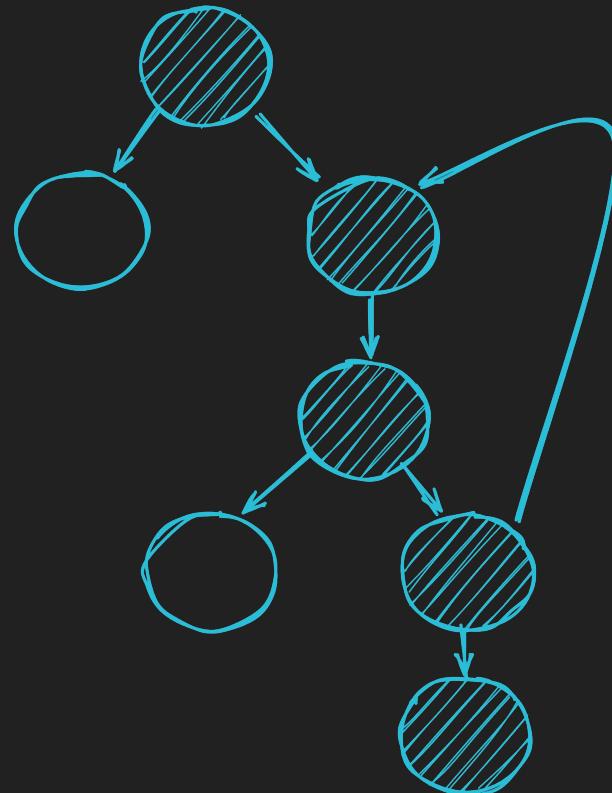
Record when nodes are covered



Control Flow: Basic Blocks

Decompose a function into a
control-flow graph

Record when nodes are covered

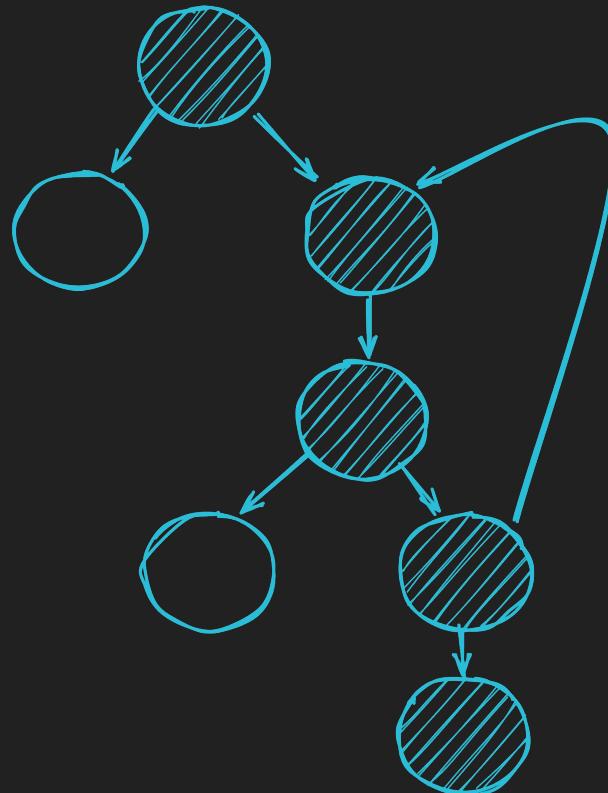


Control Flow: Basic Blocks

Decompose a function into a
control-flow graph

Record when **nodes** are covered

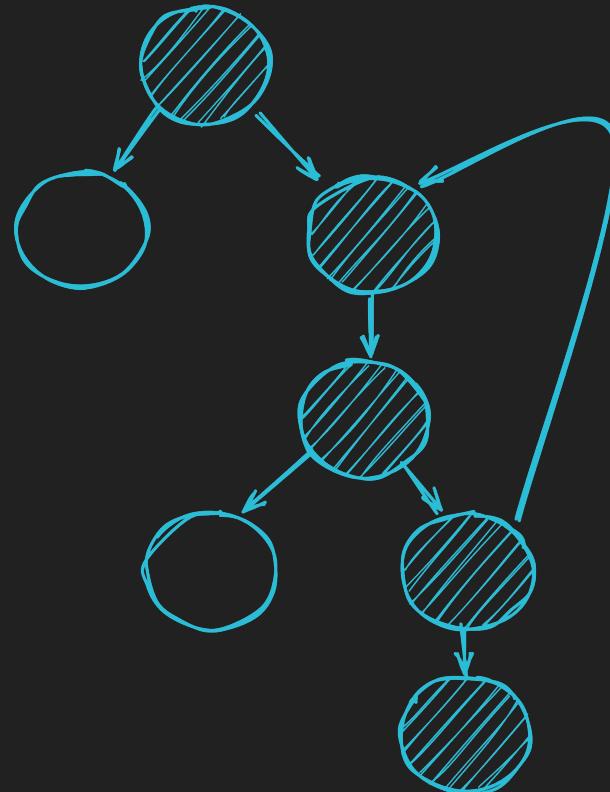
What's the problem?



Control Flow: Edges

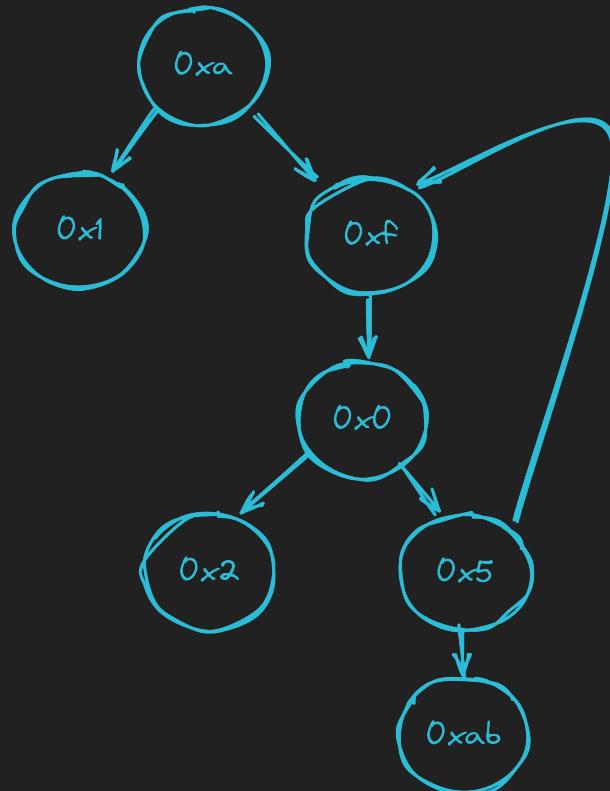
Decompose a function into a
control-flow graph

Record when ~~nodes~~ edges are covered



Control Flow: Edges

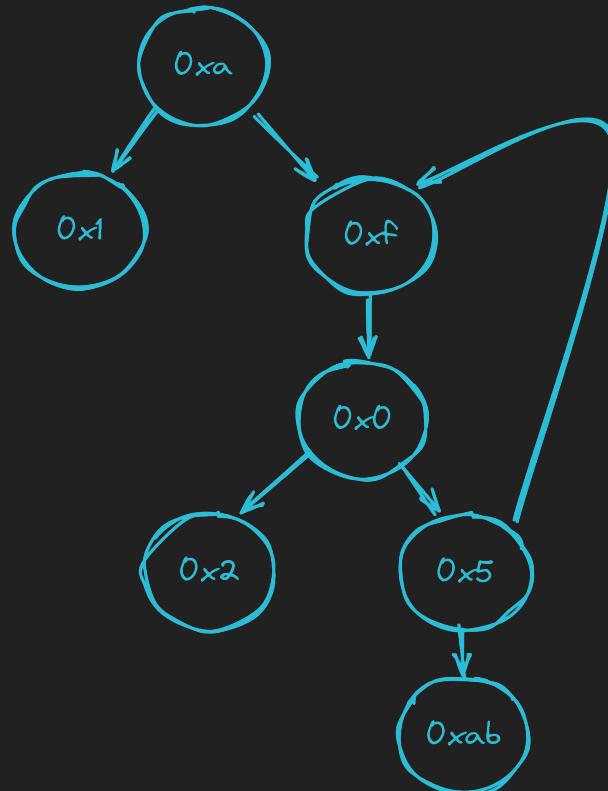
Label nodes (at compile time)



Control Flow: Edges

At start of each block (at runtime):

1. Edge ID = Prev block \wedge Curr block
2. Prev block = Right-shift Curr block

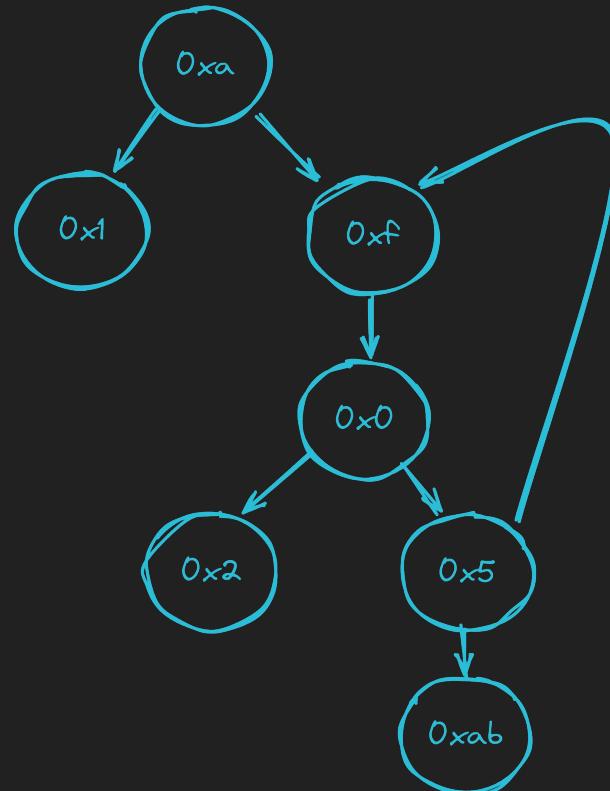


Control Flow: Edges

At start of each block (at runtime):

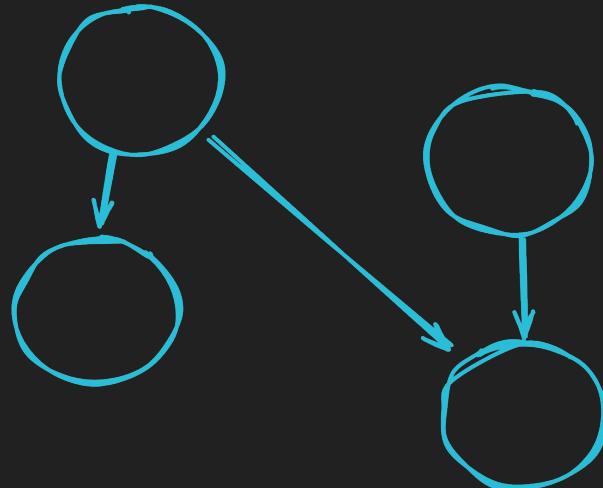
1. Edge ID = Prev block \wedge Curr block
2. Prev block = Right-shift Curr block

What's the problem?



Control Flow: “Better” Edges

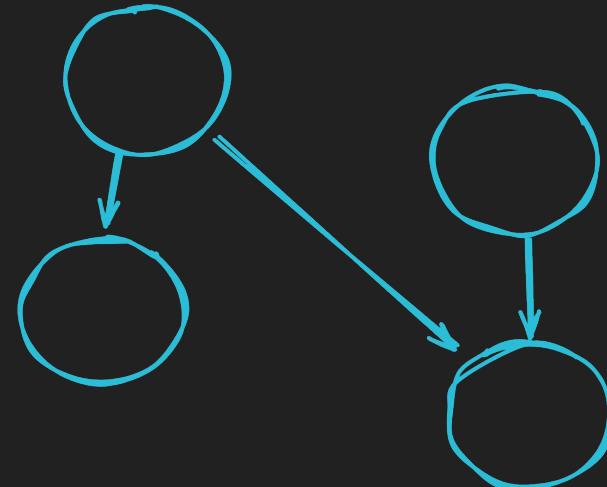
Transform the CFG and split **critical edges**



Control Flow: “Better” Edges

Transform the CFG and split **critical edges**

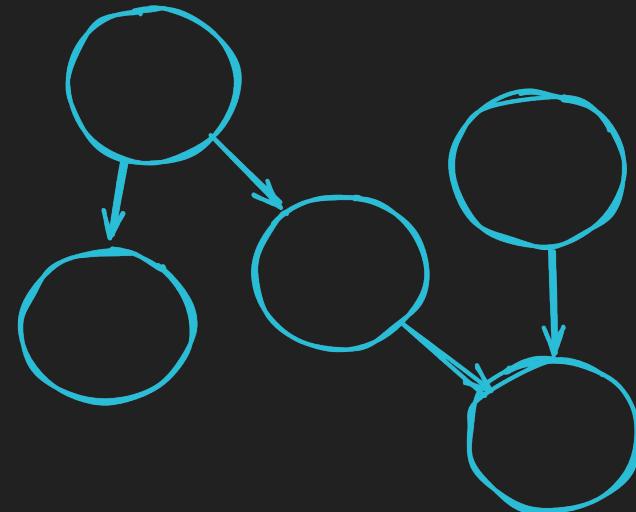
An edge whose destination has multiple predecessors and source has multiple successors



Control Flow: “Better” Edges

Transform the CFG and split **critical edges**

An edge whose destination has multiple predecessors and source has multiple successors



Insert a “dummy” block. Now, block coverage => edge coverage

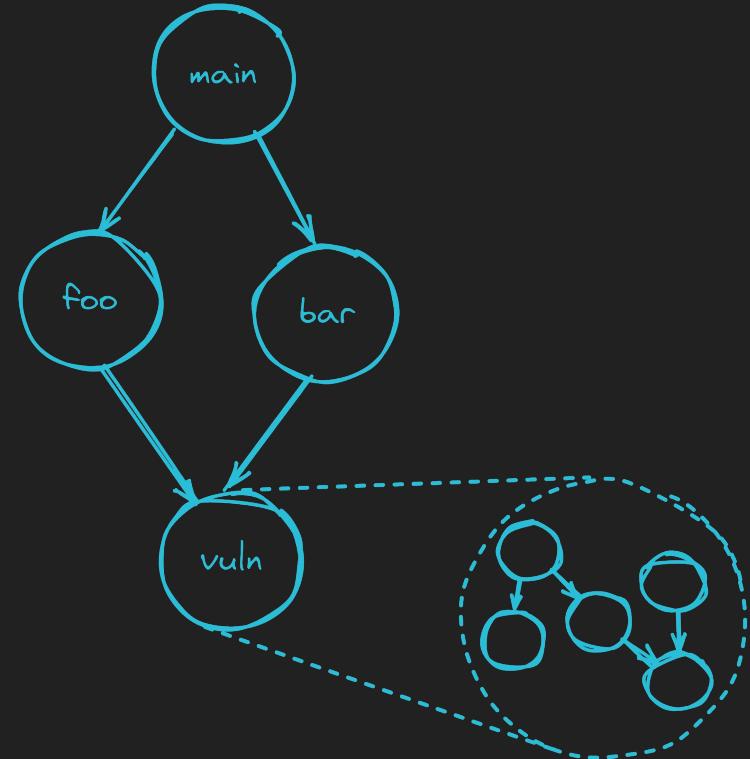
What else can we do with control flow?



Context Sensitivity

Consider the calling context

I.e., the chain of function calls leading to current location



Context Sensitivity

Label nodes and functions (at compile time)

At function call and return (at runtime):

1. Call ctx = Call ctx \wedge Function ID

At start of each block (at runtime):

1. Edge ID = Prev block \wedge Curr block \wedge Call ctx
2. Prev block = Right-shift Curr block

Context Sensitivity, Issues

Return of collisions

- Requires increasing coverage map size => slowdown

“Queue explosion”

- Retain useless seeds

Predictive Context Sensitivity

Function cloning

- Turn a context-insensitive analysis to a context-sensitive analysis
- No more collisions!

Predictive Context-sensitive Fuzzing

Pietro Borrello^{*}, Andrea Fioraldi[†], Daniele Cono D'Elia^{*},
Davide Balzarotti[‡], Leonardo Querzoni[‡] and Cristiano Giuffrida[‡]

^{*}Sapienza University of Rome

[†]EURECOM

[‡]Vrije Universiteit Amsterdam

{borrello, delta, querzoni}@diag.uniroma1.it, {fioraldi, balzarotti}@eurecom.fr, giuffrida@cs.vu.nl

Abstract. Coverage-guided fuzzers expose bugs by progressively increasing code coverage until it reaches pre-defined boundaries. Code coverage is currently the most effective and popular exploration feedback. For several bugs, though, also how contexts matter in a buggy program location may matter: for those, only tracking a coverage evolution over time and fuzzers to overlook interesting program states. Unfortunately, context-sensitive coverage tracking comes with an inherent state explosion problem. Existing attempts to solve context-sensitive coverage fuzzer struggle with it, experiencing non-trivial issues for precision (due to coverage collisions) and performance (due to context tracking and queue/map explosion).

In this paper, we show that a much more effective approach to context-sensitive fuzzing is possible. First, we propose function cloning as a basic technique to enable context-sensitive to enable precise (i.e., collision-free) context-sensitive coverage tracking. Then, to tame the state explosion problem, we argue to account for context evolution by tracking multiple contexts selected as promising. We propose a prediction scheme to identify one pool of such contexts: we analyze the data-flow diversity of each instruction and its value to call it “promising” if the fuzzer can effectively refine its value if the other sites do not.

Our work shows that, by applying function cloning to program regions that we profit to benefit from context-sensitivity, we can overcome the state explosion problem, and even improving fuzzing effectiveness. On the Fuzzelium suite, our approach largely outperforms state-of-the-art coverage-guided fuzzing embeddings, unveiling more and different bugs without increasing the number of values to call “promising”. On these heavily tested subjects, we also found 8 enduring security issues in 5 of them, with 6 CVE identifiers issued.

I. INTRODUCTION

Fuzz testing (or fuzzing for short) techniques earned a prominent place in the software security research landscape over the last decade. Their efficacy in generating unexpected or invalid inputs that make a program crash helps developers catch bugs faster than they can find them themselves. As an example, their deployment at scale in the OSS-Fuzz [2] initiative has led so far to the discovery of over 30 000 bugs in the daily testing of hundreds of open-source projects.

Network and Distributed System Security (NDSS) Symposium 2024
26 February - 1 March 2024, San Diego, CA, USA
ISBN: 1-89118-50-3
<https://doi.org/10.4727/ndss.2024.24113>
www.ndss-symposium.org

Edge coverage and other *function-local* metrics track and summarize program execution for its effects on entities (e.g., code blocks, memory locations, and file handles) [1]. A limitation of this strategy is that it may lead a fuzzer to overlook internal program states for which also *how* an entity is related matters. In program analysis, this concept goes under the name of *context-sensitivity* and is seen many applications, such as refining the precision of pointer analyses [11] and developing compiler optimizations [12].

ANGORA [11] showcases the benefits of context-sensitivity for fuzzing by augmenting edge coverage with *calling-context* information, which captures the sequence of active function calls that lead to a specific location in a program's function [13]. In principle, such a *fully context-sensitive* approach can differentiate the coverage of each test case in a fine-grained manner and lead to the discovery of more bugs [1], [10].

Predictive Context Sensitivity

Can't clone everything

- Use static analysis to inform context-sensitivity
- Favor call sites that see a higher diversity of incoming **data flow** in function arguments
- Use points-to analysis to determine diversity

Predictive Context-sensitive Fuzzing

Pietro Borrello*, Andrea Fioraldi¹, Daniele Cono D'Elia*

Davide Balzarotti¹, Leonardo Querzoni¹ and Cristiano Giuffrida²

¹Sapienza University of Rome

²EURECOM

{borrello, delta, querzoni}@diag.uniroma1.it, {fioraldi, balzarotti}@eurecom.fr, giuffrida@cs.vu.nl

Abstract. Coverage-guided fuzzers expose bugs by progressively increasing test cases until they cover all program branches. Code coverage is currently the most popular and popular exploration feedback. For several bugs, though, also *hot corners* in a buggy program location may matter: for those, only tracking a coverage evolution of fuzzer and fuzzers to overlook interesting program states. Unfortunately, context-sensitive coverage tracking comes with an inherent state explosion problem. Existing attempts to solve context-sensitive coverage in fuzzers struggle with it, experiencing non-trivial issues for precision (due to coverage collisions) and performance (due to context tracking and queue/map explosion).

In this paper, we show that a much more effective approach to context-sensitive fuzzing is possible. First, we propose function cloning as a basic technique to enable predictive to enable precise (i.e., collision-free) context-sensitive coverage tracking. Then, to tame the state explosion problem, we argue to account for context sensitivity in the fuzzer's contexts selected as promising. We propose a prediction scheme to identify one pool of such contexts: we analyze the data-flow diversity of each incoming value at a call site. Finally, the fuzzer automatically refines clusters of the latter if the latter sees incoming abstract objects that it uses at other sites do not.

Our work shows that, by applying function cloning to program regions that we predict to benefit from context-sensitivity, we can overcome the state explosion problem, and even improving fuzzing effectiveness. On the FuzzBench suite, our approach largely outperforms the state-of-the-art coverage-guided fuzzing embeddings, unveiling more and different bugs without increasing the number of values in the fuzzer's contexts. On these heavily tested subjects, we also found 8 enduring security issues in 5 of them, with 6 CVE identifiers issued.

I. INTRODUCTION

Fuzz testing (or fuzzing for short) techniques earned a prominent place in the software security research landscape over the last decade. Their efficacy in generating unexpected or invalid inputs that make a program crash helps developers catch bugs faster than they can find them themselves. As an example, their deployment at scale in the OSS-Fuzz [2] initiative has led so far to the discovery of over 30 000 bugs in the daily testing of hundreds of open-source projects.

Network and Distributed System Security (NDSS) Symposium 2024
26 February - 1 March 2024, San Diego, CA, USA
ISBN: 1-89118-93-5
<https://doi.org/10.4727/ndss.2024.24113>
www.ndss-symposium.org

Edge coverage and other *function-local* metrics track and summarize program execution for its effects on entities (e.g., code blocks, memory locations, pointers) [1].

A limitation of this strategy is that it may lead a fuzzer to overlook internal program states for which also *how* an entity is related matters. In pointer analysis, this concept goes under the name of *context-sensitivity* and is seen many applications, such as refining the precision of pointer analyses [11] and developing compiler optimizations [12].

ANGORA [11] showcases the benefits of context-sensitivity for fuzzing by augmenting edge coverage with *calling-context* information, which captures the sequence of active function calls that have led to a specific state in a program's function [13]. In principle, such a *fully context-sensitive* approach can differentiate the coverage of each test case in a fine-grained manner and lead to the discovery of more bugs [1], [10].

**There's that “data
flow” thing again...**

Data Flow Analysis

Process of collecting information about the ways variables are defined and used in a program

In compilers:

- Enables optimizations

In testing:

- Useful technique for measuring coverage

Defining Data Flow Coverage

Data Flow Analysis Techniques for Test Data Selection

Sandra Rapps* and Elaine J. Weyuker

Department of Computer Science, Courant Institute of Mathematical Sciences,
New York University, 251 Mercer Street, N.Y., N.Y. 10012

*also, YOURDN Inc., 1133 Ave. of the Americas, N.Y., N.Y. 10036

Abstract

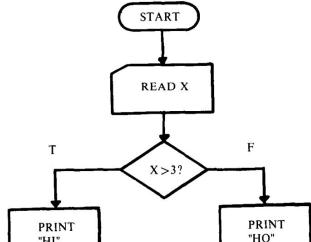
This paper examines a family of program test data selection criteria derived from data flow analysis techniques similar to those used in compiler optimization. It is argued that currently used path selection criteria which examine only the control flow of a program are inadequate. Our procedure associates with each point in a program at which a variable is defined, those points at which the value is used. Several related path criteria, which differ in the number of these associations needed to adequately test the program, are defined and compared.

Introduction

Program testing is the most commonly used method for demonstrating that a program actually accomplishes its intended purpose. The testing procedure consists of selecting elements from the program's input domain, executing the program on these test cases, and comparing the actual output with the expected output (in this discussion, we assume the existence of an "oracle", that is, some method to correctly determine the expected output). While exhaustive testing of all possible input values would provide the most complete picture of a program's performance, the size of the input domain is usually too large for this to be feasible. Instead, the usual procedure is to select a relatively small subset of the input domain which is, in some sense, representative of the entire input domain. An evaluation of the performance of the program on this test data is then used to predict its performance in general. Ideally, the test data should be chosen so that executing the program on this set will uncover all errors, thus guaranteeing that any program which produces correct results for the test data will produce correct results for any data in the input domain. However, discovering such a perfect set of test data is a difficult, if not impossible task [1,2]. In practice, test data is selected to give the tester a feeling of confidence that most errors will be discovered, without actually guaranteeing that the tested and debugged program is correct. This feeling of confidence is

select paths through the program whose elements fulfill the chosen criterion, and then to find the input data which would cause each of the chosen paths to be selected.

Using path selection criteria as test data selection criteria has a distinct weakness. Consider the strongest path selection criterion which requires that *all* program paths, p_1, p_2, \dots , be selected. This effectively partitions the input domain D into a set of classes $D = \cup D[i]$, such that for every $x \in D$, $x \in D[i]$ if and only if executing the program with input x causes path p_i to be traversed. Then, $i \in T = \{i_1, i_2, \dots\}$, where $i_j \in D[j]$, would seem to be a reasonably rigorous test of the program. However, this still does not guarantee program correctness. If one of the $D[i]$ is not revealing [2], that is for some $x_1 \in D[i]$ the program works correctly, but for some other $x_2 \in D[i]$ the program is incorrect, then if x_1 is selected as i_j the error will not be discovered. In figure 1 we see an example of this.



based on the dataflow coverage criteria. We have adapted these dataflow coverage definitions to define realistic dataflow coverage measures for C programs. A coverage measure associates a value with a set of tests for a given program. This value indicates the completeness of the set of tests for that program. We define the following dataflow coverage measures for C programs based on Rapps and Weyuker's⁷ definitions: block, decision, c-use, p-use, all-uses, path, and du-path.

Precisely defining these concepts for the C language requires some care, but the basic ideas can be illustrated by the example in Figure 1. We define the measures to be intraprocedural, so they apply equally well to individual procedures (functions), sets of procedures, or whole programs.

Block. The simplest example of a coverage measure is basic block coverage. The body of a C procedure may be considered as a sequence of basic blocks. These are portions of code that nor-

Figure 1. Sum.c computes the sum and product of numbers from 0 to N.

```
#include <stdio.h>
main()
{
    int n, i, k, sum, prod;
    printf("Enter an integer 0 for +, 1 for -: ");
    scanf("%d", &n, &k);
    sum = 0;
    prod = 1;
    i = 1;
    while (i <= n)
    {
        sum += i; // Uses of variable i
        prod *= i;
        i++;
    }
    if(k == 0)
        printf("n = %d, sum = %d", n, sum);
    if(k == 1)
        printf("n = %d, prod = %d", n, prod);
}
```



Figure 2. A hierarchy of control and dataflow coverage measures.

gram behavior, presumably due to one or more faults in the code.)

Figure 2 suggests an ordering of the coverage criteria. In this hierarchy, block coverage is weaker than decision coverage, which in turn is dominated by p-use coverage. C-use coverage dominates both block and decision coverage but is independent of p-use coverage; both c-use and

**Data-flow coverage
is the tracking of
def-use chains
executed at runtime**

Def-Use Chain Coverage

Def site: Variable allocation site (static or dynamic)

Use site: Variable access (read and/or write)

Def-use chain: Path between a def and use site

Def-Use Chain Coverage

Def site: Variable allocation site (static or dynamic)

Use site: Variable access (read and/or write)

Def-use chain: Path between a def and use site

Need an efficient implementation

A FEW
MOMENTS
LATER

datAFLow

1. Embed def-site IDs into objects
2. Reduce data-flow tracking to a metadata management problem
3. Now, def-site IDs are the metadata to retrieve at a use site

132

DatAFLow: Toward a Data-Flow-Guided Fuzzer

ADRIAN HERRERA, Australian National University, Australia
MATTHIAS PAYER, École Polytechnique Fédérale de Lausanne, Switzerland
ANTONY L. HOSKING, Australian National University, Australia

Coverage-guided greybox fuzzers rely on control-flow coverage feedback to explore a target program and uncover bugs. Compared to control-flow coverage, data-flow coverage offers a more fine-grained approximation of program behavior. Data-flow coverage captures behaviors not visible as control flow and should intuitively discover more (or different) bugs. Despite this advantage, fuzzers guided by data-flow coverage have received relatively little attention. This is due to the computational cost of data-flow analysis and its complexities (e.g., taint analysis, symbolic execution). Unfortunately, these more accurate analyses incur a high run-time overhead, impacting fuzzer throughput. Lightweight data-flow alternatives to control-flow fuzzing remain unexplored.

We present datAFLow, a greybox fuzzer guided by lightweight data-flow profiling. We also establish a framework for reasoning about data-flow coverage, allowing the computational cost of exploration to be balanced with precision. Using this framework, we extensively evaluate datAFLow across different precisions, comparing it against state-of-the-art fuzzers guided by control flow, taint analysis, and data flow.

Our results suggest that the ubiquity of control-flow-guided fuzzers is well-founded. The high run-time costs of data-flow-guided fuzzing (~10x higher than control-flow-guided fuzzing) significantly reduces fuzzer throughput. In contrast, datAFLow finds more bugs than AFL++ and AFL*, and finds many more undiscovered bugs than state-of-the-art control-flow-guided fuzzers (notably, AFL++) failed to find. This was because data-flow coverage revealed states in the target not visible under control-flow coverage. Thus, we encourage the community to continue exploring lightweight data-flow profiling; specifically, to lower run-time costs and to combine this profiling with control-flow coverage to maximize bug-finding potential.

CCS Concepts: • Software and its engineering → Software testing and debugging; Software maintenance tools; Empirical software validation;

Additional Key Words and Phrases: Fuzzing, data flow, coverage

ACM Reference format:

Adrian Herrera, Mathias Payer, and Antony L. Hosking. 2023. DatAFLow: Toward a Data-Flow-Guided Fuzzer. *ACM Trans. Softw. Eng. Methodol.* 32, 5, Article 132 (July 2023), 31 pages.
<https://doi.org/10.1145/3587156>

A. Herrera is also with Defence Science and Technology Group Next Generation Technologies Fund (Cyber) program, the DSTO Future Research Project Advanced Program Analysis for Software Vulnerability Discovery and Mitigation, SNSF PCIGP2, 186971, and ERC H2020 SG 850668.

Authors' addresses: A. Herrera and A. L. Hosking, The Australian National University, Canberra, ACT 2600, Australia; email: [adrian.herrera, antony.hosking]@anu.edu.au; M. Payer, École Polytechnique Fédérale de Lausanne, 1015 Lausanne, Switzerland; email: matthias.payer@epfl.ch

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, consider applying for permission from the copyright holder(s).
© 2023 Copyright held by the owner/author(s).
1049-311X/2023/07-ART132
<https://doi.org/10.1145/3587156>

**Can we combine
control + data flow?**

DDFuzz

- Incorporate **data dependency graph** (DDG) in coverage
- DDG represents data dependencies between instructions
- XOR into edge coverage

Fuzzing with Data Dependency Information

Alessandro Mantovani
EURECOM
mantovan@eurecom.fr

Andrea Fioraldi
EURECOM
fioraldi@eurecom.fr

Davide Balzarotti
EURECOM
balzaro@eurecom.fr

Abstract—Recent advances in fuzz testing have introduced several forms of feedback mechanisms, motivated by the fact that for a large range of programs and libraries, edge-coverage alone is insufficient to reveal complicated bugs. Inspired by this line of research, we proposed existing program representations to be leveraged for such heterogeneous purposes. In this paper, we propose to extend the notion of the structure and adaptability to the context of fuzz testing. In particular, we believe that data dependency graphs (DDGs) represent a good candidate for this task, as the set of information embedded by this data structure is potentially useful for fuzzing. We demonstrate this by introducing associations of defense pairs that would be the result for a traditional fuzzer to trigger. Since some portions of the dependency graph overlap with the control flow of the program, it is possible to reduce the additional instrumentation to cover only “interesting” data-flow paths of the program, those that help the fuzzer to visit the code in a distinct way compared to standard branch coverage.

To test these observations, in this paper we propose DDFuzz, a new approach that rewards the fuzzer not only with code coverage information, but also when new edges in the data dependency graph are hit. Our results show that the adoption of data dependency instrumentation in coverage-guided fuzzing is a promising solution that can help to discover bugs that are not easily found by standard coverage approaches. This is demonstrated by the 72 different vulnerabilities that our data-dependency driven approach can identify when executed on 38 target programs from three different datasets.

1. Introduction

In a society that makes software applications the central core of many every-day activities is essential to make such software as secure as possible before it is released to the public. This has led to a large amount of research focused on the development of increasingly sophisticated techniques to discover vulnerabilities such as static software testing [36], [60], [77], symbolic execution [61], [62], [71] and dynamic analysis [73].

In the context of dynamic analysis, researchers have proposed many techniques to increase the coverage of a certain metric produced in the software under testing. One of the possible metrics is *path coverage*, which refers to all independent paths present in a program. For example, in software testing, the community has focused on path coverage for tests generation [64], [70] with the goal of automatically producing inputs that can reach nested code locations. The main limitation of path coverage is that any non trivial application can contain a huge, and potentially infinite, number of paths [48], [67] thus making this approach a poor choice for a large set of programs.

Because of this, other dynamic testing techniques, like *branch coverage*, mostly rely on simpler forms of coverage, such as the popular *edge coverage*.

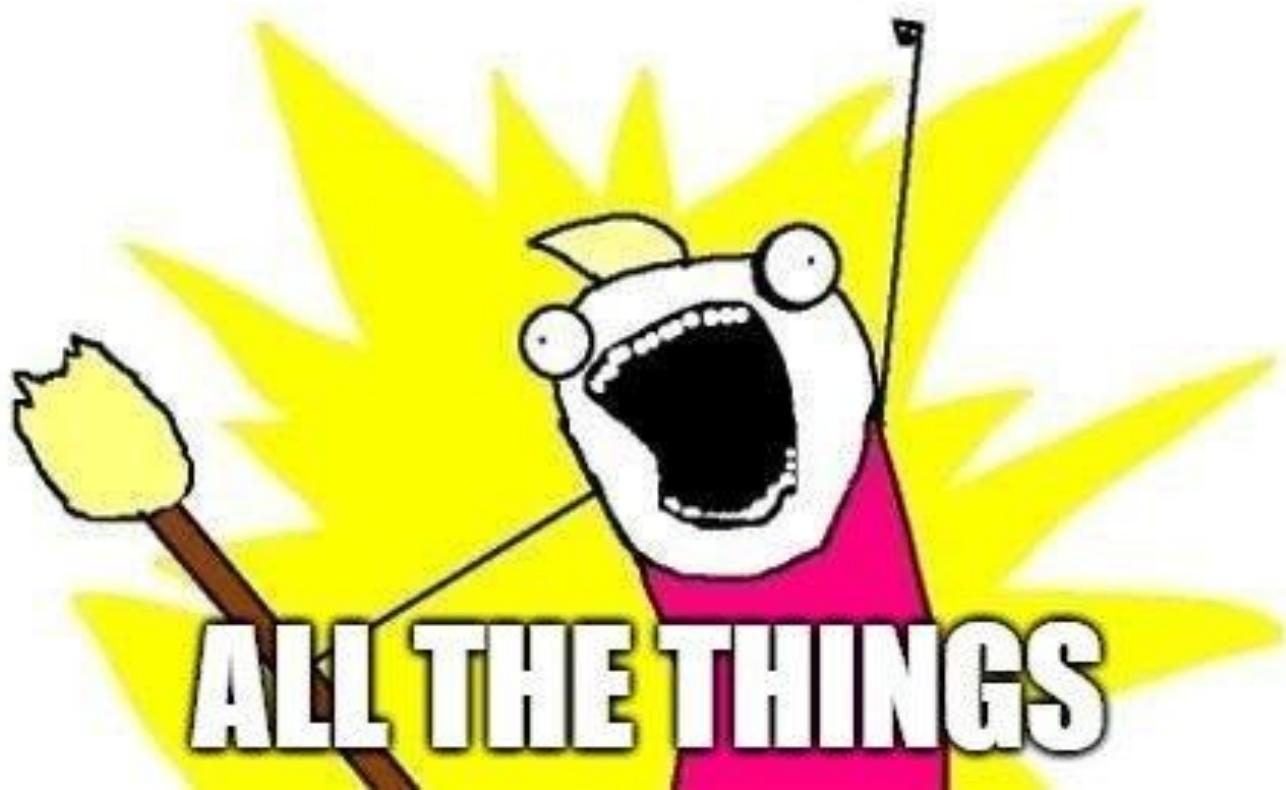
Fuzz testing or *fuzzing* is a dynamic analysis technique that consists of feeding a program with different variants, with the goal of revealing a flaw in its underlying code. Originally proposed in the early ‘90 [52], fuzz testing has evolved enormously over the past decade thanks to numerous efforts that have been made to improve its individual components. For instance, researchers have tried to enhance fuzzers by designing better seed schedulers [21], [66] or by introducing novel mutation engines [31], [32], [49]. Today, *edge coverage* testing is probably nowadays the most effective approach employed for automatic vulnerability discovery and for this reason public research on this topic has advanced extremely fast as confirmed by the fact that in 2010 alone more than 120 papers were published on this topic [51].

As a necessary condition for a fuzzer to discover a vulnerability is the ability to generate an input that reaches the location of the flaw in the target program, *code coverage* has become one of the most common metrics used by the fuzzing community. One of the main goals of a fuzzing campaign. In this context, *edge coverage* became the de-facto standard to measure code coverage. According to this paradigm, all basic blocks in the control flow graph (CFG) of the binary are instrumented to reward the fuzzer with a feedback when new edges connecting two basic blocks are hit. In this way, the fuzzer can keep track of the new discovered parts of the program point and mutate the generated inputs so that they evolve towards the exploration of nested portions of code. This gave origin to what is commonly referred as *coverage-guided fuzzing*.

Consequently, one of the main goals of the fuzzing researchers is to develop new techniques to increase edge coverage, which has led current state-of-the-art fuzzers to be very effective at visiting difficult-to-trigger code locations. For instance, RedQueen [11] can generate inputs that satisfy complicated conditions and the ones that implement a function with some magic bytes. However, high edge coverage alone does not always translate in a better ability to discover bugs, and therefore the fuzzing community had also explored other approaches based on alternative coverage metrics. In this direction, Parnesan [55] relies on the code instruments instantiated with `AddressSanitizer` [1] and `UndefinedBehaviorSanitizer` [1] to build a

**So what actually
works?**

FUZZ



Key Findings

Speed matters

- Dumb + fast > smart + slow

Different coverage metrics find different bugs

- This occurs even when coverage of one metric is less than another

In most programs, control flow subsumes data flow

Key Questions

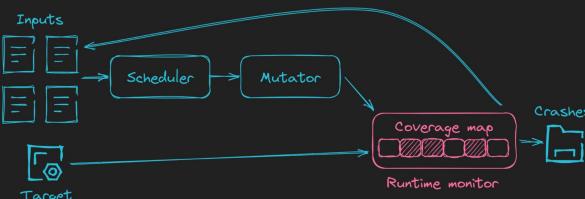
- What other ways can we approximate a program's state space?
- Can we perform an initial (static?) analysis of the target to guide what coverage metric to use?
- Ensemble techniques?

The Hitchhiker's Guide to Fuzzer Coverage Metrics



Me as a grad student
(sometimes)

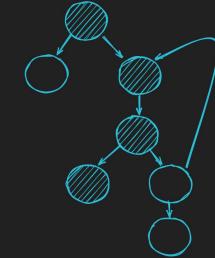
Fuzzing 101



Control Flow: Basic Blocks

Decompose a function into a control-flow graph

Record when nodes are covered



Predictive Context Sensitivity

Can't clone everything

- Use static analysis to inform context-sensitivity
- Favor call sites that see a higher diversity of incoming **data flow** in function arguments
- Use points-to analysis to determine diversity



datAFLow

1. Embed def-site IDs into objects
2. Reduce data-flow tracking to a metadata management problem
3. Now, def-site IDs are the metadata to retrieve at a use site



DDFuzz

- Incorporate **data dependency graph** (DDG) in coverage
- DDG represents data dependencies between instructions
- XOR into edge coverage



Key Findings

Speed matters

- Dumb + fast > smart + slow

Different coverage metrics find different bugs

- This occurs even when coverage of one metric is less than another



In most programs, control flow subsumes data flow

Key Questions

- What other ways can we approximate a program's state space?
- Can we perform an initial (static?) analysis of the target to guide what coverage metric to use?
- Ensemble techniques?